

STAT 545 HW 4

2026-02-09

1

Consider a regression model with $p < n$:

$$y = X\beta + \epsilon$$

where $y \in R^n$, $X \in R^{n \times p}$ and $\epsilon \sim N(0, \sigma^2 I_n)$ and $\text{rank}(X) = p$. Recall that the ridge regression estimate is given by

$$\hat{\beta}_{ridge} = (X^T X + \lambda I)^{-1} X^T y$$

for given $\lambda > 0$.

A.

Let the SVD of X be given by $X = UDV^T$. Show that the fitted value is

$$\hat{y} = X\hat{\beta}_{ridge} = \sum_{j=1}^p u_j \frac{d_j^2}{d_j^2 + \lambda} u_j^T y$$

where u_j are the columns of U and $D = \text{diag}(d_j)$.

First, since it is SVD, we know that $U^T U = V^T V = I_p$. We know the columns of U and V form an orthonormal bases for the row and column space of X respectively.

We first want to calculate $\hat{\beta}_{ridge}$.

We know from chapter 2 notes, we can find

$$X^T X = (V D U^T)(U D V^T) = V D^2 V^T$$

Then, we can get it in a nice form like in the slides

$$(X^T X + \lambda I) = V D^2 V^T + \lambda(V V^T) = V(D^2 + \lambda I)V^T$$

Then, simple inverse (Is possible because we added lambda values along diagonal, making it positive definite). Note that the inverse of an orthonormal matrix is its own transpose.

$$(X^T X + \lambda I)^{-1} = (V(D^2 + \lambda I)V^T)^{-1} = V(D^2 + \lambda I)^{-1}V^T$$

We can now plug everything in and find

$$\begin{aligned}\hat{\beta}_{ridge} &= V(D^2 + \lambda I)^{-1}V^T(VDU^T)y \\ &= V(D^2 + \lambda I)^{-1}DU^T y\end{aligned}$$

Multiply by X now.

$$\begin{aligned}\hat{y} &= X\hat{\beta}_{ridge} \\ &= (UDV^T)V(D^2 + \lambda I)^{-1}DU^T y \\ &= UD(D^2 + \lambda I)^{-1}DU^T y \\ &= \sum_{j=1}^p u_j \frac{d_j^2}{d_j^2 + \lambda} u_j^T y\end{aligned}$$

Each u_j is a column vector multiplied by the value of the j-th diagonal value in D (a diagonal matrix, $p \times p$), so we can treat $D(D^2 + \lambda I)^{-1}D$ as $\frac{d_j^2}{d_j^2 + \lambda}$ and split the calculation of \hat{y} as a sum of the operations on the j columns.

B.

When the fitted value \hat{y} has a form $\hat{y} = My$ where M is a matrix that does not depend on y then the “degrees of freedom” is given by the trace of the matrix M. Derive an expression (a scalar quantity) for the degrees of freedom of ridge regression.

We know

$$\begin{aligned}\hat{y} &= X\hat{\beta}_{ridge} \\ &= (UDV^T)V(D^2 + \lambda I)^{-1}DU^T y \\ &= UD(D^2 + \lambda I)^{-1}DU^T y \\ \implies M &= UD(D^2 + \lambda I)^{-1}DU^T\end{aligned}$$

Then

$$\begin{aligned}tr(M) &= tr(UD(D^2 + \lambda I)^{-1}DU^T) \\ &= tr(D(D^2 + \lambda I)^{-1}DU^T U) \\ &= tr(D(D^2 + \lambda I)^{-1}D) \\ &= \sum_{j=1}^p \frac{d_j^2}{d_j^2 + \lambda}\end{aligned}$$

Given that the degree of freedom is equal to the trace, we find that $df = tr(M) = \sum_{j=1}^p \frac{d_j^2}{d_j^2 + \lambda}$.

C.

Take $n = 70$, $p = 100$. Simulate X and ϵ from i.i.d. standard normals. Take β to be a vector of all ones and set $y = X\beta + \epsilon$. Take $\lambda = 1$ and verify that (a) that the ridge estimate is well-defined and (b) both ways of calculating the fitted value above (via matrix inverse and via SVD) indeed result in the same \hat{y} in R. The commands svd() for calculating SVD may be useful.

First we set it all up.

```

set.seed(13)
n <- 70
p <- 100
X <- matrix(rnorm(n * p), nrow = n, ncol = p)
epsilon <- rnorm(n)
beta <- rep(1, p)
y <- X %*% beta + epsilon
lambda <- 1

```

Direct computation way

```

beta_ridge <- solve(t(X) %*% X + lambda * diag(p)) %*% t(X) %*% y
y_hat_inv <- X %*% beta_ridge

```

```

# Check values
beta_ridge

```

```

## [1] 1.253267328
## [2] 0.824068812
## [3] 0.522741183
## [4] 0.978814764
## [5] 0.466439308
## [6] 1.250124085
## [7] 0.321746645
## [8] 1.240042943
## [9] 0.319374310
## [10] 0.979175475
## [11] 0.665154813
## [12] 1.899374467
## [13] 0.414776278
## [14] 1.486085861
## [15] 0.769178241
## [16] 0.682495639
## [17] 0.312515779
## [18] 0.608304003
## [19] 1.090411287
## [20] 0.647268079
## [21] 1.935008401
## [22] 0.774677535
## [23] 0.503690748
## [24] 1.576099086
## [25] 0.883043617
## [26] 0.436607704
## [27] 0.902727906
## [28] 1.402553778
## [29] 1.090050344
## [30] 0.874637508
## [31] 0.990564358
## [32] 0.376301286
## [33] 1.066534515

```

```

## [34,] -0.060808858
## [35,]  0.133728141
## [36,]  0.591757043
## [37,]  0.938083854
## [38,]  1.153184984
## [39,] -0.003031669
## [40,]  0.857557656
## [41,] -0.300039862
## [42,]  0.670482295
## [43,]  0.863936516
## [44,]  0.093572456
## [45,]  0.371717743
## [46,]  0.833470563
## [47,]  0.807508827
## [48,]  1.036199830
## [49,]  1.844007927
## [50,]  1.116752435
## [51,]  0.984484274
## [52,]  0.682835724
## [53,]  0.775713528
## [54,]  0.850923473
## [55,]  0.479990672
## [56,]  0.144934484
## [57,]  0.533261128
## [58,]  1.636666867
## [59,]  0.200729498
## [60,]  0.422847735
## [61,]  1.058942156
## [62,]  0.282593915
## [63,]  0.948986634
## [64,]  0.619825939
## [65,]  0.626653348
## [66,]  0.317341709
## [67,]  0.473644646
## [68,]  0.536327016
## [69,] -0.104662967
## [70,] -0.050013051
## [71,]  1.021453146
## [72,]  0.298307614
## [73,]  0.776355498
## [74,]  0.663582287
## [75,]  0.720490266
## [76,]  0.626932288
## [77,]  1.114055982
## [78,]  1.449874503
## [79,]  0.683645985
## [80,] -0.102013800
## [81,]  1.494683561
## [82,]  0.307068367
## [83,]  0.468604824
## [84,]  0.787521556
## [85,]  1.284092059
## [86,]  0.948890748
## [87,]  0.407425831

```

```
## [88,] 0.630615860
## [89,] 1.184312322
## [90,] 0.453168164
## [91,] 0.447722957
## [92,] 0.720061814
## [93,] 0.651741818
## [94,] 0.790570056
## [95,] 0.888219252
## [96,] 0.726008138
## [97,] 0.616305138
## [98,] 0.928586003
## [99,] 0.734734715
## [100,] 0.075364795
```

```
y_hat_inv
```

```
## [,1]
## [1,] 11.6930595
## [2,] -2.2828304
## [3,] 8.3389511
## [4,] -17.3525478
## [5,] 0.7857293
## [6,] -3.0377488
## [7,] 9.8427591
## [8,] 8.9135213
## [9,] 2.0486252
## [10,] -1.2649932
## [11,] 11.4920223
## [12,] -15.9795011
## [13,] -5.0582121
## [14,] -3.0245648
## [15,] -1.3977781
## [16,] -3.7033322
## [17,] 8.0426574
## [18,] -23.7889261
## [19,] -4.6927534
## [20,] 16.1170924
## [21,] 4.6581993
## [22,] 6.2699274
## [23,] -8.3129443
## [24,] 10.8168377
## [25,] -14.1247429
## [26,] 2.8640045
## [27,] -20.9264941
## [28,] -3.3688354
## [29,] 2.1126005
## [30,] -20.6679985
## [31,] 6.6283992
## [32,] 17.8736022
## [33,] -21.2346204
## [34,] 6.2344835
## [35,] -11.2643976
## [36,] 5.6457907
## [37,] 13.6744408
```

```

## [38,] -11.9604539
## [39,] 1.3115843
## [40,] -16.1456468
## [41,] -14.2352921
## [42,] 10.0227490
## [43,] -7.8865110
## [44,] 7.9407564
## [45,] -16.8534783
## [46,] 22.0800211
## [47,] -17.7969466
## [48,] 13.1852053
## [49,] -7.5489433
## [50,] 8.7510592
## [51,] 11.5745645
## [52,] 9.4745742
## [53,] -5.5360794
## [54,] 3.6792100
## [55,] -0.7659997
## [56,] 4.4279193
## [57,] 4.9919759
## [58,] 17.4007937
## [59,] 6.8553628
## [60,] 2.3306133
## [61,] -6.4441001
## [62,] 9.8295794
## [63,] 11.8660865
## [64,] 6.0520626
## [65,] -20.8839999
## [66,] 8.2596429
## [67,] 1.3960509
## [68,] 4.9397570
## [69,] 3.0433631
## [70,] -16.3029813

```

The SVD way

We find β_{ridge} and then compute the estimate \hat{y} .

```

SVD_X <- svd(X)
SVD_U <- SVD_X$u
SVD_D <- SVD_X$d

# We use the formula from part A.
SVD_y_hat <- matrix(0, nrow = n, ncol = 1)
for (j in 1:length(SVD_D)) {
  uj <- SVD_U[, j, drop = FALSE]
  d_diagonals <- (SVD_D[j]^2) / (SVD_D[j]^2 + lambda)
  UTY <- t(uj) %*% y
  SVD_y_hat <- SVD_y_hat + uj %*% (d_diagonals * UTY)
}

# Check our results
print(SVD_y_hat)

```

```

## [,1]
## [1,] 11.6930595
## [2,] -2.2828304
## [3,] 8.3389511
## [4,] -17.3525478
## [5,] 0.7857293
## [6,] -3.0377488
## [7,] 9.8427591
## [8,] 8.9135213
## [9,] 2.0486252
## [10,] -1.2649932
## [11,] 11.4920223
## [12,] -15.9795011
## [13,] -5.0582121
## [14,] -3.0245648
## [15,] -1.3977781
## [16,] -3.7033322
## [17,] 8.0426574
## [18,] -23.7889261
## [19,] -4.6927534
## [20,] 16.1170924
## [21,] 4.6581993
## [22,] 6.2699274
## [23,] -8.3129443
## [24,] 10.8168377
## [25,] -14.1247429
## [26,] 2.8640045
## [27,] -20.9264941
## [28,] -3.3688354
## [29,] 2.1126005
## [30,] -20.6679985
## [31,] 6.6283992
## [32,] 17.8736022
## [33,] -21.2346204
## [34,] 6.2344835
## [35,] -11.2643976
## [36,] 5.6457907
## [37,] 13.6744408
## [38,] -11.9604539
## [39,] 1.3115843
## [40,] -16.1456468
## [41,] -14.2352921
## [42,] 10.0227490
## [43,] -7.8865110
## [44,] 7.9407564
## [45,] -16.8534783
## [46,] 22.0800211
## [47,] -17.7969466
## [48,] 13.1852053
## [49,] -7.5489433
## [50,] 8.7510592
## [51,] 11.5745645
## [52,] 9.4745742
## [53,] -5.5360794

```

```

## [54,]  3.6792100
## [55,] -0.7659997
## [56,]  4.4279193
## [57,]  4.9919759
## [58,] 17.4007937
## [59,]  6.8553628
## [60,]  2.3306133
## [61,] -6.4441001
## [62,]  9.8295794
## [63,] 11.8660865
## [64,]  6.0520626
## [65,] -20.8839999
## [66,]  8.2596429
## [67,]  1.3960509
## [68,]  4.9397570
## [69,]  3.0433631
## [70,] -16.3029813

```

Check results

```
max(abs(y_hat_inv - SVD_y_hat))
```

```
## [1] 5.631051e-13
```

The two predictions that we got through the two different methods actually give us very similar results. They are nearly exactly the same.

We have a totally good prediction for \hat{y} even though in theory, it is not possible to get the inverse of $X^T X$ for regular LS approach but with our ridge regression, we can still calculate some sort of inverse that is close and get predicted values.