

# STAT 545 HW 5

Bowen Zheng

2026-02-18

## Problem 1

### A.

We can try to see how long it takes each of the parts to compute to find out.

Inverting D would be just taking reciprocal of p items, so it is  $O(p)$  complexity.

Finding  $X^T X$  would be multiplying a  $p \times n$  and  $n \times p$  matrices. This would take  $O(p^2 n)$ .

Adding two matrices  $X^T X + D^{-1}$  would only take  $p \times p$  computations, so  $O(p^2)$ .

Finally, the most complex step, the last one is inverting the  $p \times p$  matrix so  $O(p^3)$ .

The biggest / dominating computational complexity is in this last step and thus our computational complexity for directly computing  $\Sigma$  is  $O(p^3)$ .

### B.

We choose the obvious matrices to match up with the Woodbury matrix identity:

$A = D^{-1}$  (a  $p \times p$  diagonal matrix)

$U = X^T$  (a  $p \times n$  matrix)

$B = I_n$  (an  $n \times n$  identity matrix)

$V = X$  (an  $n \times p$  matrix)

Then we find

$$\Sigma = D - DX^T(I_n + XDX^T)^{-1}XD$$

Again, we can look at each step again to figure out the complexity.

Multiplying  $DX^T$  or  $XD$  is simply one computation per thing so we get  $O(np)$ .

Computing the  $XDX^T$  will be  $O(np)$ , then multiplying the  $n \times p$  by  $p \times n$  will be  $O(n^2 p)$ .

Adding  $I_n + XDX^T$  will be adding two  $n \times n$  so  $O(n^2)$ .

Inverting  $I_n + XDX^T$  will be inverting  $n \times n$  so  $O(n^3)$ .

Matrix multiplication of the parts together will be  $n \times p$  by  $p \times n$  or  $p \times n$  by  $n \times p$  so  $O(np^2)$  or  $O(n^2 p)$ .

Since  $p > n$ , our computational complexity  $O(np^2)$  should mean it is less complex than in part (A.) where it was  $O(p^3)$ .

## C.

### First the distribution

In step 3, we know  $(XDX^T + I_n)w = (y - v)$ , then  $w = (XDX^T + I_n)^{-1}(y - v)$ , where  $v = Xu + \delta$ . Since  $E[u] = 0$  and  $E[\delta] = 0$ , it follows that  $E[v] = 0$ . Then,  $E[w] = (XDX^T + I_n)^{-1}(y - 0)$  and with step 4,  $\theta = u + DX^T w$ .

We can then find the expectation  $E[\theta] = E[u] + DX^T E[w] = 0 + DX^T(XDX^T + I_n)^{-1}y$ .

Now, We find that  $X^T(XDX^T + I_n) = X^T XDX^T + X^T = (D^{-1} + X^T X)DX^T$ . We multiply by  $(D^{-1} + X^T X)^{-1}$  on the left and multiply by  $(XDX^T + I_n)^{-1}$  on the right, to find

$$(D^{-1} + X^T X)^{-1} X^T (XDX^T + I_n) (XDX^T + I_n)^{-1} = (D^{-1} + X^T X)^{-1} (D^{-1} + X^T X) D X^T (XDX^T + I_n)^{-1}$$

$$(D^{-1} + X^T X)^{-1} X^T = D X^T (XDX^T + I_n)^{-1}$$

Note that this is a very useful identity we will use multiple times. Another one we will provide now:

$$\begin{aligned}\Sigma &= (D^{-1} + X^T X)^{-1} \\ \Sigma^{-1} &= D^{-1} + X^T X \\ I_p &= \Sigma D^{-1} + \Sigma X^T X \\ (I_p - \Sigma X^T X) &= \Sigma D^{-1}\end{aligned}$$

Rewrite  $(D^{-1} + X^T X)^{-1} X^T = (X^T X + D^{-1})^{-1} X^T = \Sigma$ . Then, from earlier

$$\begin{aligned}E[\theta] &= 0 + D X^T (XDX^T + I_n)^{-1} y \\ &= (X^T X + D^{-1})^{-1} X^T y \\ &= \Sigma X^T y \\ &= \mu\end{aligned}$$

Now we need to solve for the covariance matrix. Again, we use step 4 identity:  $\theta = u + DX^T w$ . We substitute in our  $w$ , our formula we just solved for the  $\Sigma$ , remember that  $u\delta$  independent and find

$$\begin{aligned}\theta &= u + D X^T (XDX^T + I_n)^{-1} (y - v) \\ &= u + D X^T (XDX^T + I_n)^{-1} (y - Xu - \delta) \\ &= (I_p - D X^T (XDX^T + I_n)^{-1} X) u + (D X^T (XDX^T + I_n)^{-1})(y - \delta) \\ &= (I_p - (D^{-1} + X^T X)^{-1} X^T X) u + (D^{-1} + X^T X)^{-1} X^T (y - \delta) \\ &= (I_p - \Sigma X^T X) u + \Sigma X^T (y - \delta) \\ \implies Var(\theta) &= (I_p - \Sigma X^T X) Var(u) (I_p - \Sigma X^T X)^T + (\Sigma X^T) I_n (\Sigma X^T)^T \\ &= (\Sigma D^{-1}) D (\Sigma D^{-1})^T + (\Sigma X^T) I_n (\Sigma X^T)^T \\ &= \Sigma D^{-1} \Sigma + \Sigma X^T X \Sigma \\ &= \Sigma (D^{-1} + X^T X) \Sigma \\ &= \Sigma \Sigma^{-1} \Sigma \\ &= \Sigma\end{aligned}$$

$\theta$  is a linear combination of normal distributions with mean  $\mu$  and variance  $\Sigma$ . Thus,  $\theta \sim N(\mu, \Sigma)$ .

## Computational complexity

Step 1: We first sample p and then n, so  $O(p + n)$  complexity there.

Step 2: We have to do matrix vector multiplication,  $n \times p$  by  $p \times 1$  so  $O(np)$  and then add the two vectors  $O(n)$ .

Step 3: To solve for w, we have to take the inverse of the  $n \times n$  ( $XDX^T + I_n$ ) and then multiply it by  $(y - v)$ , which is  $O(n^2p)$  and  $O(n^3)$ . Since  $p > n$ , then  $O(n^2p)$  is more.

Step 4: This final step just involves addition and multiplication of matrices. Adding is  $O(n)$  and multiplication is  $O(np)$ .

The most dominant term is  $O(n^2p)$  coming from step 3 that takes  $O(n^2p + n^3)$  computational complexity.

## D.

Take  $n = 50$ ,  $p = 100$ . Simulate  $X$  and  $\epsilon$  from i.i.d. standard normals. Take  $\beta$  to be a vector of all ones and set  $y = X\beta + \epsilon$ . Take  $D = I$  and implement the sampling schemes of parts (a) , (b) and (c) to generate 1,000 samples each. Report the computational times. You may use the Cholesky decomposition function in R.

```
# Data Generation
set.seed(218)
n <- 50
p <- 100
n_samples <- 1000

X <- matrix(rnorm(n * p), n, p)
epsilon <- rnorm(n)
beta_true <- rep(1, p)
y <- X %*% beta_true + epsilon
D_diag <- rep(1, p)
sqrt_D_diag <- sqrt(D_diag)
D <- diag(D_diag)

# Part A: Doing steps then sampling
time_a <- system.time({
  D_inv <- solve(D)
  Precision <- t(X) %*% X + D_inv
  Sigma <- solve(Precision)
  mu <- Sigma %*% t(X) %*% y
  L <- t(chol(Sigma))
  samples_a <- replicate(n_samples, mu + L %*% rnorm(p))
})

# Part B: Woodbury for Sigma
time_b <- system.time({
  In <- diag(n)
  D_inv <- solve(D)
  inner_inv <- solve(In + X %*% D_inv %*% t(X))
  Sigma_w <- D_inv - D_inv %*% t(X) %*% inner_inv %*% X %*% D_inv
  mu_w <- Sigma_w %*% t(X) %*% y
  L_w <- t(chol(Sigma_w))
```

```

samples_b <- replicate(n_samples, mu_w + L_w %*% rnorm(p))
})
# Part C: Algorithm
time_c <- system.time({
  In <- diag(n)
  inner_mat <- X %*% D %*% t(X) + In

  samples_c <- replicate(n_samples, {
    u <- sqrt_D_diag * rnorm(p)
    delta <- rnorm(n)
    v <- X %*% u + delta
    w <- solve(inner_mat, y-v)
    theta <- u + t(X) %*% w
    theta
  })
})

# Reporting Times
cat("Part (A) Direct: ", time_a["elapsed"], "s\n")

```

```
## Part (A) Direct: 0.04 s
```

```
cat("Part (B) Woodbury: ", time_b["elapsed"], "s\n")
```

```
## Part (B) Woodbury: 0.02 s
```

```
cat("Part (C) Algorithm: ", time_c["elapsed"], "s\n")
```

```
## Part (C) Algorithm: 0.07 s
```

Each run returns a different amount of time taken. In theory, part B and C should take the least time. However, it seems that from fastest to slowest is B A C. Of course, it swings wildly, but after running the same chunk of code many many times, this seems to be the correct pattern. In fact, the order in which I generate / time the sampling procedure seems to matter, but if I generate A B C, then B A C is the order I get generally. This is not the case if I generate the samples in a different order. Perhaps it would be more consistent if the sample size was not so small.